

Artificial Intelligence: Through the Lens of Poker

Max Chiswick

Table of contents

Welcome	8
I Introduction	9
1 Why Poker?	10
2 Poker Camp	11
3 Critical Thinking and Problem Solving	12
4 Games	13
5 Basic Strategy	14
6 Ethical Considerations	15
II Coding Foundations	16
7 Python	17
8 AI Agents	18
9 Building AI Agents	19
III Math Foundations	20
10 Algebra	21
11 Probability	22
12 Combinatorics	23
13 Expected Value	24
14 Bayes' Rule	25

15 Statistics	26
16 Monte Carlo Methods	27
17 Data Analysis	28
IV Knowledge	29
18 Logic	30
19 Knowledge Representation	31
20 Rationality	32
21 Psychology and Mindset	33
V Decision Making Under Uncertainty	34
22 Probabilistic Thinking	35
23 Decision Theory	36
24 Rule-Based Systems	37
25 Risk	38
26 Regret	39
27 Multi-armed Bandits	40
VI Game Theory	41
28 Nash Equilibrium	42
29 Game Theory Optimal (GTO)	43
30 Mixed Strategies	44
VII Game Trees	45
31 Perfect Information	46

32 Search and Minimax	47
33 Imperfect Information Games	48
VIII Solving Toy Games	49
34 Analytical Solutions	50
35 Normal Form	51
36 Optimization	52
37 Sequence Form	53
IX Counterfactual Regret Minimization	54
38 CFR Algorithm	55
38.1 Counterfactual	55
38.2 Regret	55
38.3 Minimization	55
38.4 What is CFR?	56
38.5 Algorithm Overview	56
38.6 Simplified CFR Example: Regret Updates in RPS	57
38.7 Simplified CFR Example: Counterfactual Values	61
38.8 Core CFR Algorithm	62
38.8.1 Regret and Strategies	63
38.8.2 Iterating through the Tree	65
38.9 More Exercises	66
39 CFR Interactive	69
40 Solving RPS and Poker with CFR	70
40.0.1 RPS Regret Matching Experiment	70
40.1 Data: Talk Paper Scissors	70
41 CFR Proof	72
42 CFR Algorithm Improvements	74
42.1 CFR+	74
42.2 Discounted CFR	74
43 Monte Carlo CFR	75

44	Vector CFR	76
X	Abstracting Large Games	77
45	Game Size	78
46	Card Abstraction	79
47	Bet Abstraction	80
48	Agent Evaluation	81
XI	Poker Solvers	82
49	How Solvers Work	83
50	Using Solvers	84
51	Generalizing Solver Outputs	85
52	Studying Populations	86
53	Solver Limitations	87
54	Advanced Strategy	88
55	Tournaments	89
XII	AI Math Foundations	90
56	Calculus	91
57	Linear Algebra	92
58	Information Theory	93
XIII	AI Foundations	94
59	Machine Learning	95
60	Deep Learning	96

61 Reinforcement Learning	97
XIVRecent Game AI Advances	98
62 Types of Games	99
63 AlphaGo and AlphaZero	100
64 Multi-agent Systems	101
XV State of the Art Poker AI	102
65 Deep CFR	103
66 Top Poker Agents	104
67 Variance Reduction	105
68 Human vs. AI	106
69 New Research	107
XVILLMs	108
70 Transformers	109
71 OthelloGPT	110
72 PokerGPT	111
73 Interpretability	112
XVIDopponent Modeling	113
74 Best Response	114
75 Exploitative Strategies	115
XVIAI Risks and Safety	116
76 Ethics and Short-term Risks	117

77 Alignment and Long-term Risks	118
XIX The River	119
78 Trading	120
79 Prediction Markets	121
80 Other Betting	122
XX Project Ideas	123
81 Projects	124

Welcome

A text by [Poker Camp](#). Coming soon.

Part I

Introduction

1 Why Poker?

Hello

2 Poker Camp

Hello

3 Critical Thinking and Problem Solving

Hello

4 Games

Hello

5 Basic Strategy

Hello

6 Ethical Considerations

Hello

Part II

Coding Foundations

7 Python

Hello

8 AI Agents

Hello

9 Building AI Agents

Hello

Part III

Math Foundations

10 Algebra

Hello

11 Probability

Hello

12 Combinatorics

Hello

13 Expected Value

Hello

14 Bayes' Rule

Hello

15 Statistics

Hello

16 Monte Carlo Methods

Hello

17 Data Analysis

Hello

Part IV

Knowledge

18 Logic

Hello

19 Knowledge Representation

Hello

20 Rationality

Hello

21 Psychology and Mindset

Hello

Part V

Decision Making Under Uncertainty

22 Probabilistic Thinking

Hello

23 Decision Theory

Hello

24 Rule-Based Systems

Hello

25 Risk

Hello

26 Regret

Hello

$$\text{Regret}(a) = u(a, s_{-i}) - u(s_i, s_{-i})$$

where:

a : the action being evaluated

s_i : the strategy actually played

s_{-i} : the opponent's strategy

$u(\cdot, \cdot)$: the utility function

27 Multi-armed Bandits

Hello

Part VI

Game Theory

28 Nash Equilibrium

Hello

29 Game Theory Optimal (GTO)

Hello

30 Mixed Strategies

Hello

Part VII

Game Trees

31 Perfect Information

Hello

32 Search and Minimax

Hello

33 Imperfect Information Games

Hello

Part VIII

Solving Toy Games

34 Analytical Solutions

Hello

35 Normal Form

Hello

36 Optimization

Hello

37 Sequence Form

Hello

Part IX

Counterfactual Regret Minimization

38 CFR Algorithm

The most popular method for iteratively solving poker games is the Counterfactual Regret Minimization (CFR) algorithm. CFR was developed in 2007 at the University of Alberta. It converges to Nash equilibrium in two player zero-sum games.

38.1 Counterfactual

Counterfactual means “relating to or expressing what has not happened or is not the case”. For example, if in reality I didn’t bring an umbrella and got wet in the rain, I could say counterfactually, “If I had brought an umbrella, I wouldn’t have gotten wet.”

Actual event: I didn’t bring an umbrella, and I got wet in the rain

Counterfactual event: If I had brought an umbrella, I wouldn’t have gotten wet

38.2 Regret

Regret we previously touched on in [Chapter 26](#).

In brief, it’s a way to assign a value to the difference between a made decision and an optimal decision. For example, if you choose to play a slot machine that returns a value of 5 rather than the best machine that returns a value of 10, then your regret would be $10 - 5 = 5$.

38.3 Minimization

Minimization refers to minimizing the difference between the made decision and the optimal decision. Playing the optimal slot machine, i.e. the one that returns a value of 10, would minimize the regret to 0.

38.4 What is CFR?

CFR is a self-play algorithm that learns by playing against itself repeatedly. It starts with a uniform random strategy (each action at each decision point is equally likely) and iterates on these strategies each round to nudge closer to the game theory optimal Nash equilibrium strategy.

Each action at each information set in the game has an associated regret value that is calculated based on how that action performs compared to how the overall strategy at that info set performs. Positive regret means that the action did better than the overall strategy and negative means worse.

The regrets get updated each iteration and upon returning to that node for the next iteration, the strategy is calculated proportionally to the positive regrets, meaning that actions with higher regret that performed well in the past will get played more often.

At the end, the average of all strategies played converges to the equilibrium strategy.

The strategy is computed offline and then can be used in play – it’s a fixed, opponent-agnostic strategy that can’t be beaten in the long-run, but also doesn’t take advantage of opponent weaknesses.

See (Johanson 2015) for more detail.

38.5 Algorithm Overview

Each information set maintains a **strategy** and **regret** tabular counter for each action. These accumulate the sum of all strategies and the sum of all regrets.

In a game like Rock Paper Scissors, there is effectively only one info set, so only one table for **strategy** over each action (Rock, Paper, Scissors) and one table for **regret** over each action (Rock, Paper, Scissors).

Regrets are linked to strategies through a policy called *regret matching*, which selects strategies proportionally to their positive regrets.

After using regret matching and after many iterations, we minimize expected regret by using the *average strategy* at the end, which is the strategy that converges to equilibrium.

If two players were training against each other using regret matching, they would converge to the Nash Equilibrium of 1/3 for each action using the average strategy in Rock Paper Scissors.

38.6 Simplified CFR Example: Regret Updates in RPS

Let's see how that works in Rock Paper Scissors.

In general, we define regret as:

$$\text{Regret} = u(\text{Alternative Strategy}) - u(\text{Current Strategy})$$

$$\text{Regret}(a) = u(a, a_{\text{opp}}) - u(a_{\text{played}}, a_{\text{opp}})$$

where:

a : the action being evaluated for regret

a_{played} : the action actually played

a_{opp} : the opponent's action

$u(\cdot, \cdot)$: the utility function

We prefer alternative actions with high regret *and* wish to minimize our overall regret.

Wait, isn't that counterintuitive?

Well, by definition we want to play actions with higher regret because they perform better. But how does this minimize our overall regret?

Because as we play more of these actions, then our current strategy and

We play Rock and opponent plays Paper $\Rightarrow u(\text{rock}, \text{paper}) = -1$

$$\text{Regret}(\text{paper}) = u(\text{paper}, \text{paper}) - u(\text{rock}, \text{paper}) = 0 - (-1) = 1$$

$$\text{Regret}(\text{rock}) = u(\text{rock}, \text{paper}) - u(\text{rock}, \text{paper}) = -1 - (-1) = 0$$

$$\text{Regret}(\text{scissors}) = u(\text{scissors}, \text{paper}) - u(\text{rock}, \text{paper}) = 1 - (-1) = 2$$

We play Scissors and opponent plays Paper $\Rightarrow u(\text{scissors}, \text{paper}) = 1$

$$\text{Regret}(\text{paper}) = u(\text{paper}, \text{paper}) - u(\text{scissors}, \text{paper}) = 0 - 1 = -1$$

$$\text{Regret}(\text{rock}) = u(\text{rock}, \text{paper}) - u(\text{scissors}, \text{paper}) = -1 - 1 = -2$$

$$\text{Regret}(\text{scissors}) = u(\text{scissors}, \text{paper}) - u(\text{scissors}, \text{paper}) = 1 - 1 = 0$$

We play Paper and opponent plays Paper $\Rightarrow u(\text{paper}, \text{paper}) = 0$

$$\text{Regret}(\text{paper}) = u(\text{paper}, \text{paper}) - u(\text{paper}, \text{paper}) = 0 - 0 = 0$$

$$\text{Regret}(\text{rock}) = u(\text{rock}, \text{paper}) - u(\text{paper}, \text{paper}) = -1 - 0 = -1$$

$$\text{Regret}(\text{scissors}) = u(\text{scissors}, \text{paper}) - u(\text{paper}, \text{paper}) = 1 - 0 = 1$$

To generalize:

- The action played always gets a regret of 0 since the “alternative” is really just that same action
- When we play a tying action, the alternative losing action gets a regret of -1 and the alternative winning action gets a regret of $+1$
- When we play a winning action, the alternative tying action gets a regret of -1 and the alternative losing action gets a regret of -2
- When we play a losing action, the alternative winning action gets a regret of $+2$ and the alternative tying action gets a regret of $+1$

After each play, we accumulate regrets for each of the 3 actions.

We decide our strategy probability distribution using regret matching, which means playing a strategy that normalizes over the *positive* accumulated regrets, i.e. playing in proportion to the positive regrets.

Here’s an example from (Neller and Lanctot 2013). Assume that the regret counters start at 0 for each action.

- Game 1: Choose Rock and opponent chooses Paper
 - Lose 1
 - Rock: Regret 0
 - Paper: Regret 1
 - Scissors: Regret 2

- Next Action: Proportional


$$\begin{pmatrix} \text{Rock} & 0/3 = 0 \\ \text{Paper} & 1/3 = 0.333 \\ \text{Scissors} & 2/3 = 0.667 \end{pmatrix}$$

- Game 2: Choose Scissors (probability $2/3$) and opponent chooses Rock
 - Lose 1
 - Rock: Regret 1
 - Paper: Regret 2
 - Scissors: Regret 0
- Cumulative regrets:
 - Rock: 1
 - Paper: 3
 - Scissors: 2

- Next Action: Proportional

$$\begin{pmatrix} \text{Rock} & 1/6 = 0.167 \\ \text{Paper} & 3/6 = 0.500 \\ \text{Scissors} & 2/6 = 0.333 \end{pmatrix}$$

38.7 Simplified CFR Example: Counterfactual Values



Here we have a very simple game tree showing a Player 1 node at the end of the game. For simplicity, we assume that this is the first iteration of the algorithm so each action is assigned a uniform 1/3 probability.

The counterfactual reach probability of information state I_i , $\pi_{-i}(I_i)$, is the probability of reaching I_i with strategy profile σ_{-i} except that, we treat current player i actions to reach the state as having probability 1. In all situations we refer to as “counterfactual”, one treats the computation as if player i ’s strategy was modified to have intentionally played to information set I_i . Put another way, we exclude the probabilities that factually came into player i ’s play from the computation.

Regular expected value: We can compute the regular expected value for each action of P1:

$$\mathbb{E}(a) = 0.33 * (4) = 1.33$$

$$\mathbb{E}(b) = 0.33 * (7) = 2.33$$

$$\mathbb{E}(c) = 0.33 * (-10) = -3.33$$

Counterfactual value: The counterfactual value includes the reach probability $\sigma_{-i}(I_i) = 0.38$ of the opponent and chance playing to this node.

$$CV(I_i, a) = \pi_{-i}(I_i) \pi^{I_i \rightarrow a}(h, z) u_i(z)$$

38.8 Core CFR Algorithm

High quantity of definitions (conventions primarily sourced from 2015 paper [Solving Heads-up Limit Texas Hold'em](#)):

P : Set of players

H : Game state, represented as history of actions from start of game

Tabular storing strategies and regrets at each info set

Regrets based on action values compared to node EV, which is based on counterfactual values

Regret minimization, usually regret matching, to get new strategies

Average strategy converges to Nash equilibrium

As we think about solving larger games, we start to look at iterative algorithms.

38.8.1 Regret and Strategies

A strategy at an info set is a probability distribution over each possible action.

Regret is a measure of how much each strategy at an info set is preferred and is used as a way to update strategies.

For a given P1 strategy and P2 strategy, a player has regret when they take an action at an info set that was not the highest-EV action at that info set.

Regret Exercise

What is the regret for each action?

Action	Regret
A	
B	
C	

Solution

Action	Regret
A	4
B	2
C	0

Expected Value Exercise

If taking a uniform strategy at this node (i.e. $\frac{1}{3}$ for each action), then what is the expected value of the node?

Solution

$$\mathbb{E} = \frac{1}{3} * 1 + \frac{1}{3} * 3 + \frac{1}{3} * 5 = 0.33 + 1 + 1.67 = 3$$

Poker Regret Exercise

In poker games, the regret for each action is defined as the value for that action minus the expected value of the node. Give the regret values for each action under this definition.

Solution

Action	Value	Poker Regret
A	1	-2
B	3	0
C	5	2

At a node in a poker game, the player prefers actions with higher regrets by this definition.

Regret matching definitions:

- a is actions
- σ is strategy
- t is time
- i is player
- R is cumulative regret

$$\sigma_i^t(a) = \begin{cases} \frac{\max(R_i^t(a), 0)}{\sum_{a' \in A} \max(R_i^t(a'), 0)} & \text{if } \sum_{a' \in A} \max(R_i^t(a'), 0) > 0 \\ \frac{1}{|A|} & \text{otherwise} \end{cases}$$

This is showing that we take the cumulative regret for an action divided by the cumulative regrets for all actions (normalizing) and then play that strategy for this action on the next iteration.

If all cumulative regrets are ≤ 0 then we use the uniform distribution.

If cumulative regrets are positive, but are < 0 for a specific action, then we use 0 for that action.

In code:

```
def get_strategy(self):
    #First find the normalizing sum
    normalizing_sum = 0
    for a in range(NUM_ACTIONS):
        if self.regret_sum[a] > 0:
            self.strategy[a] = self.regret_sum[a]
        else:
            self.strategy[a] = 0
    normalizing_sum += self.strategy[a]

    #Then normalize each action
    for a in range(NUM_ACTIONS):
        if normalizing_sum > 0:
```



```

        self.strategy[a] /= normalizing_sum
    else:
        self.strategy[a] = 1.0/NUM_ACTIONS
        self.strategy_sum[a] += self.strategy[a]

    return self.strategy

```

38.8.2 Iterating through the Tree

The core feature of the iterative algorithms is self-play by traversing the game tree over all **infosets** and tracking the strategies and regrets at each.

From above, we know how to find the strategy and regret in the simple Rock Paper Scissors environment.

In poker:

- Strategies are determined the same as above, through regret matching from the previous **regret** values at the specific information set for each action
- CFR definitions:
 - a is actions
 - I is infoset
 - σ is strategy
 - t is time
 - i is player
 - R is cumulative regret
 - z is a terminal node
 - u is utility (payoffs)
 - p is the current player who plays at this node
 - $-p$ is the the opponent player and chance
 - v is counterfactual value
- Counterfactual values are effectively the value of an information set. They are weighted by the probability of opponent and chance playing to this node (in other words, the probability of playing to this node if this player tried to do so).
 - Counterfactual value: $v^\sigma(I) = \sum_{z \in Z_I} \pi_{-p}^\sigma(z[I]) \pi^\sigma(z[I] \rightarrow z) u_p(z)$
 - $\sum_{z \in Z_I}$ is summing over all terminal histories reachable from this node
 - $\pi_{-p}^\sigma(z[I])$ is the probability of opponents and chance reaching this node

- $\pi^\sigma(z[I] \rightarrow z)$ is the probability of playing from this node to terminal history z , i.e. the weight component of the expected value
- $u_p(z)$ is the utility at terminal history z , i.e. the value component of the expected value
- Instantaneous regrets are based on action values compared to info-set EV. Each action EV then adds to its **regret** counter:
 - $r^t(I, a) = v^{\sigma^t}(I, a) - v^{\sigma^t}(I)$
- Cumulative (counterfactual) regrets are the sum of the individual regrets:
 - $R^T(I, a) = \sum_{t=1}^T r^t(I, a)$

38.9 More Exercises

Maximize Against non-Nash Fixed Opponent

How would you maximize in RPS knowing the opponent plays a fixed non-Nash strategy that you don't know?

Solution

One option is to play the equilibrium strategy until you get a significant sample on your opponent and then to exploit their strategy going forward.

Strategy Against No-Rock Opponent

What is the optimal play if your opponent can't play Rock?

Solution

Player 1/2	Paper	Scissors
Rock	(-1, 1)	(1, -1)
Paper	(0, 0)	(-1, 1)
Scissors	(1, -1)	(0, 0)

We can see that Player 1 playing Paper is dominated by Scissors, so Player 1 should never play Paper.

Player 1/2	Paper	Scissors
Rock	$(-1, 1)$	$(1, -1)$
Scissors	$(1, -1)$	$(0, 0)$

In the reduced game, we see that if Player 2 plays Paper with probability p and Scissors with probability s , then:

$$\mathbb{E}(\text{P1 R}) = -1 * p + 1 * s = -p + s \quad \mathbb{E}(\text{P1 S}) = 1 * p + 0 * s = p$$

Setting these equal, $-p + s = p \Rightarrow s = 2p$.

We also know that $s + p = 1$.

Therefore $s = 1 - p$ and $1 - p = 2p \Rightarrow 1 = 3p \Rightarrow p = 1/3$.

Therefore, $s = 1 - 1/3 = 2/3$.

For Player 2, we have $s = 2/3$ and $p = 1/3$.

For Player 1, we can solve similarly:

$$\mathbb{E}(\text{P2 P}) = 1 * r - 1 * s = r - s \quad \mathbb{E}(\text{P2 S}) = -1 * r + 0 * s = -r$$

$$r - s = -r \Rightarrow 2r = s$$

We also know that $r + s = 1$.

Therefore $s = 1 - r$ and $1 - r = 2r \Rightarrow 1 = 3r \Rightarrow r = 1/3$.

Therefore, $s = 1 - 1/3 = 2/3$.

For Player 2, we have $s = 2/3$ and $p = 1/3$.

Inserting these probabilities, we have:

Player 1/2	Paper (1/3)	Scissors (2/3)
Rock (1/3)	$(-1, 1) (1/9)$	$(1, -1) (2/9)$
Scissors (2/3)	$(1, -1) (2/9)$	$(0, 0) (4/9)$

Therefore Player 1 has payoffs of: $1/9 * -1 + 2/9 * 1 + 2/9 * 1 + 4/9 * 0 = 3/9 = 1/3$.

Therefore the player that can still play Rock has an advantage of $1/3$ at equilibrium.

Maximize Against Adapting Rock Opponent

1. Suppose that your opponent is forced to play Rock exactly

Suppose that your opponent gets a card with probability X such that $X\%$ of the time they are forced to play Rock What if the opponent is adapting to you, but 10% of the time they are forced to play Rock?

Solution

Soon

Skewed Rock Payoff

What is the equilibrium strategy if the payoff for Rock over Scissors is 2 (others stay the same)?

Solution

Soon

39 CFR Interactive

Hello

40 Solving RPS and Poker with CFR

Bet capping: Even when holding the strongest hand — a pair of aces — the strategy caps the betting less than 0.01% of the time

40.0.1 RPS Regret Matching Experiment

Here we show that regret matching converges only using the average strategy over 10,000 iterations:

The bottom shows both players converging to 1/3, while the top shows Player 1's volatile current strategies that are cycling around.

Suppose that your opponent Player 2 is playing 40% Rock, 30% Paper, and 30% Scissors. Here is a regret matching 10,000 game experiment. It shows that it takes around 1,600 games before Player 1 plays only Paper (this will vary).

We see that if there is a fixed player, regret matching converges to the best strategy.

But what if your opponent is not using a fixed strategy? We'll talk about that soon.

40.1 Data: Talk Paper Scissors

[eieio games](#) made a Rock Paper Scissors over voice game in which players call a phone number and get matched up with another player for a 3 game RPS match.

They published their 40,000 round data on X:

Overall: R 37.2%, P 35.4%, S 27.4%

Round 1: R 39.7%, P 37.6%, S 22.7%

Round 2: R 34.0%, 33.4%, 32.6%

Round 3: R 37.2%, 34.7%, 28.1%

Expected Value Against TPS Player

What is the best strategy per round against the average TPS player? What is your expected value per round and overall?

Solution

The best strategy is to always play Paper.

$$\mathbb{E}(\text{Round 1}) = 0.397 * 1 + 0.376 * 0 + 0.227 * -1 = 0.17$$

$$\mathbb{E}(\text{Round 2}) = 0.34 * 1 + 0.334 * 0 + 0.326 * -1 = 0.014$$

$$\mathbb{E}(\text{Round 3}) = 0.372 * 1 + 0.347 * 0 + 0.281 * -1 = 0.091$$

$$\mathbb{E}(\text{Round 4}) = 0.17 + 0.014 + 0.091 = 0.275$$

41 CFR Proof

Hello

Intuitions

https://www.reddit.com/r/GAMETHEORY/comments/nq7gly/doubts_about_counterfactual_regret_minim
<https://www.quora.com/>

Epsilon-Nash Equilibrium HULHE is solved

A game is said to be ultraweakly solved if, for the initial position(s), the game-theoretic value has been determined; weakly solved if, for the initial position(s), a strategy has been determined to obtain at least the game-theoretic value, for both players, under reasonable resources; and strongly solved if, for all legal positions, a strategy has been determined to obtain the game-theoretic value of the position, for both players, under reasonable resources. In an imperfect-information game, where the game-theoretic value of a position beyond the initial position is not unique, Allis's notion of "strongly solved" is not well defined. Overall, we require less than 11 TB of storage to store the regrets and 6 TB to store the average strategy during the computation, which is distributed across a cluster of computation nodes. This amount is infeasible to store in main memory, and so we store the values on each node's local disk. But: Finally, unlike with CFR, we have empirically observed that the exploitability of the players' current strategies during the computation regularly approaches zero. Therefore, we can skip the step of computing and storing the average strategy, instead using the players' current strategies as the CFR+ solution.

A strategy's "exploitability" is the maximum amount that a perfect counter-strategy could win on expectation against a strategy. A Nash equilibrium has an exploitability of zero, since it cannot be beaten by anyone on expectation, and having a lower exploitability is good. When you run CFR, the average strategy's exploitability converges towards zero, driving its worst-case loss lower and lower. Note that this is a pessimistic way to measure how good your strategy is: our best poker programs started beating the world's best human players in heads-up limit hold'em in 2008, even though our programs at that time were still massively exploitable by this worst-case measure.

In our January 2015 Science paper, we've announced that we've produced a strategy that has essentially weakly solved the game. That means that we have computed a strategy with such a low exploitability (0.000986 big blinds per game) that it would take more than a human lifetime of play, using the perfect counter-strategy, for anyone to have 95% statistical confidence that they were actually winning against it. So it's not an exactly perfect strategy, but it is so close

to perfect that the game is essentially solved, as it's now outside of any human's ability to beat it for a statistically meaningful amount by playing games against it.

42 CFR Algorithm Improvements

Hello

CFR+ variation such that regrets can't be < 0

Linear CFR such that regrets are weighted by their recency

42.1 CFR+

42.2 Discounted CFR

<https://www.cs.cmu.edu/~sandholm/cs15-888F23/> lecture 6 When is CFR+ worse? High variance play keeps getting very positive

CFR+ does exhaustive iterations over the entire game tree and uses a variant of regret matching (regret matching+) where regrets are constrained to be non-negative. Actions that have appeared poor (with less than zero regret for not having been played) will be chosen again immediately after proving useful (rather than waiting many iterations for the regret to become positive). Finally, unlike with CFR, we have empirically observed that the exploitability of the players' current strategies during the computation regularly approaches zero. Therefore, we can skip the step of computing and storing the average strategy, instead using the players' current strategies as the CFR+ solution.

What are other ways to improve CFR?

43 Monte Carlo CFR

Hello

<https://www.cs.cmu.edu/~sandholm/cs15-888F21/> lecture 10 Improve upon CFR with Monte Carlo sampling Monte Carlo CFR versions (focus on chance sampling and external sampling) MCCFR Mini-Project: MCCFR implemented <https://www.cs.cmu.edu/~sandholm/cs15-888F23/> lecture 14 <https://ai.stackexchange.com/questions/26345/how-exactly-is-monte-carlo-counterfactual-regret-minimization-with-external-samp>

Sampling methods

External: Sample chance and opponent nodes

Chance: Sample chance only

Outcome: Sample outcomes

44 Vector CFR

Hello

MCCFR worse for heads-up postflop, slows down at lower exploitability Vector alg efficient at heads-up postflop (pass down ranges, update regrets for entire range at once), evaluate showdowns $O(n)$ instead of $O(n^2)$

Part X

Abstracting Large Games

45 Game Size

Hello

46 Card Abstraction

Hello

47 Bet Abstraction

Hello

48 Agent Evaluation

Hello

Part XI

Poker Solvers

49 How Solvers Work

Hello

50 Using Solvers

Hello

51 Generalizing Solver Outputs

Hello

52 Studying Populations

Hello

53 Solver Limitations

Hello

54 Advanced Strategy

Hello

55 Tournaments

Hello

Part XII

AI Math Foundations

56 Calculus

Hello

57 Linear Algebra

Hello

58 Information Theory

Hello

Part XIII

AI Foundations

59 Machine Learning

Hello

60 Deep Learning

Hello

61 Reinforcement Learning

Hello

Part XIV

Recent Game AI Advances

62 Types of Games

Hello

63 AlphaGo and AlphaZero

Hello

64 Multi-agent Systems

Hello

Part XV

State of the Art Poker AI

65 Deep CFR

Hello

66 Top Poker Agents

Hello

67 Variance Reduction

Hello

68 Human vs. AI

Hello

69 New Research

Hello

Part XVI

LLMs

70 Transformers

Hello

71 OthelloGPT

Hello

72 PokerGPT

Hello

73 Interpretability

Hello

Part XVII

Opponent Modeling

74 Best Response

Hello

75 Exploitative Strategies

Hello

Part XVIII

AI Risks and Safety

76 Ethics and Short-term Risks

Hello

77 Alignment and Long-term Risks

Hello

Part XIX

The River

78 Trading

Hello

79 Prediction Markets

Hello

80 Other Betting

Hello

Part XX

Project Ideas

81 Projects

Hello

Johanson, Michael. 2015. “What Is an Intuitive Explanation of Counterfactual Regret Minimization?” Quora. 2015. <https://www.quora.com/What-is-an-intuitive-explanation-of-counterfactual-regret-minimization>.

Neller, Todd W., and Marc Lanctot. 2013. “An Introduction to Counterfactual Regret Minimization.” Technical Report. Gettysburg, PA: Gettysburg College. <http://modelai.gettysburg.edu/2013/cfr/cfr.pdf>.